

Author: Aditya Jain

About : Assignment 4, APS1080 (Introduction to RL)

Topic : Function Approximation and SARSA Control

```
In [2]: import matplotlib.pyplot as plt
import gym
from IPython import display as ipythondisplay
import numpy as np
import json
import pickle
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import torch
```

```
In [3]: env = gym.make('MountainCar-v0')

print('Observation Space: ', env.observation_space)
print('Action Space: ', env.action_space)

no_actions = 3
```

Observation Space: Box(-1.2000000476837158, 0.6000000238418579, (2,)), float32)

Action Space: Discrete(3)

Initialize the Neural Network

```
In [4]: input_dim = 2 # position and velocity
no_layers1 = 256
no_layers2 = 128
no_layers3 = 128
out_dim = no_actions # no of actions

model = keras.Sequential()
model.add(keras.Input(shape=(input_dim,)))
model.add(layers.Dense(no_layers1, activation="relu"))
model.add(layers.Dense(no_layers2, activation="relu"))
model.add(layers.Dense(no_layers3, activation="relu"))
model.add(layers.Dense(out_dim))

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 256)	768
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 128)	16512
dense_3 (Dense)	(None, 3)	387
Total params: 50,563		
Trainable params: 50,563		

Non-trainable params: 0

## Epsilon greedy action function

In [5]:

```
def epsilon_greedy_action(obs, epsilon, no_actions, model):
    '''chooses epsilon greedy action given a policy'''

    if np.random.rand() < epsilon:
        action = np.random.choice(no_actions, p=[1/no_actions, 1/no_actions,
    else:
        obs      = tf.expand_dims(obs, axis=0)
        pred     = model(obs)
        action = np.argmax(pred)

    return action
```

## Semi-gradient SARSA Training

In [ ]:

```
episodes      = 0
tot_step      = 0
epsilon       = 0.1
alpha         = 1e-3
optimizer     = keras.optimizers.SGD(learning_rate=alpha)
loss_fn       = keras.losses.MeanSquaredError()
gamma         = 0.9
converged     = False
conv_step     = 0
conv_thr      = 2000      # if loss not decreasing after n epsidoes, then converge
epi_cost      = []
min_loss      = 1e+6

while not converged:
    episodes += 1
    epi_rew  = 0

    cur_obs  = env.reset()
    cur_action = epsilon_greedy_action(cur_obs, epsilon, no_actions, model)
    done     = False

    while not done:
        next_obs, reward, done, info = env.step(cur_action)
        epi_rew += reward
        tot_step += 1
        if done:
            with tf.GradientTape() as tape:
                target      = reward
                prediction   = model(tf.expand_dims(cur_obs, axis=0), trainable_weights=model.trainable_weights)
                cur_estimate = prediction[0][cur_action]
                loss_value   = loss_fn(tf.expand_dims(target, axis=0), tf.expand_dims(cur_estimate, axis=0))
            else:
                next_action = epsilon_greedy_action(next_obs, epsilon, no_actions, model)
                with tf.GradientTape() as tape:
                    next_prediction = model(tf.expand_dims(next_obs, axis=0), trainable_weights=model.trainable_weights)
                    target          = reward + gamma*next_prediction[0][next_action]
                    cur_prediction  = model(tf.expand_dims(cur_obs, axis=0), trainable_weights=model.trainable_weights)
                    cur_estimate    = cur_prediction[0][cur_action]
                    loss_value      = loss_fn(tf.expand_dims(target, axis=0), tf.expand_dims(cur_estimate, axis=0))
                cur_obs      = next_obs
                cur_action    = next_action

            experiment.log_metric("Loss per timestep", loss_value, step=tot_step)
            grads = tape.gradient(loss_value, model.trainable_weights)
```

```
optimizer.apply_gradients(zip(grads, model.trainable_weights))

if loss_value < min_loss:
    min_loss = loss_value
    conv_step = 0
else:
    conv_step += 1

# if conv_step >= conv_thr:
#     converged = True
if episodes>400:
    converged = True

epi_cost.append(epi_rew)
experiment.log_metric("episode cost", epi_rew, step=episodes)

experiment.end()
```

In [7]:

```
model.save('semigradient_a4.h5')
```

WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile\_metrics` will be empty until you train or evaluate the model.

### Testing the trained controller

In [8]:

```
def control_performance(env_name, model, no_actions, trials):
    env = gym.make(env_name)
    reward_list = []
    steps_list = []
    reached_goal = 0

    for i in range(trials):
        done = False
        episode_reward = 0
        obs = env.reset()
        steps = 0

        while not done:
            obs = tf.expand_dims(obs, axis=0)
            pred = model(obs)
            action = np.argmax(pred)
            obs, reward, done, info = env.step(action)
            steps += 1
            episode_reward += reward

            if steps<200:
                steps_list.append(steps)
                reached_goal += 1

        reward_list.append(episode_reward)

    return reached_goal, steps_list
```

In [ ]: